OPERATING SYSTEMS CPU Scheduling

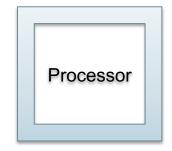
Basic Concepts

- Maximum CPU utilization obtained with multiprogramming
- Continuous Cycle:
 - one process has to wait (I/O)
 - Operating system takes the CPU away
 - Give CPU to another process
 - This pattern continues



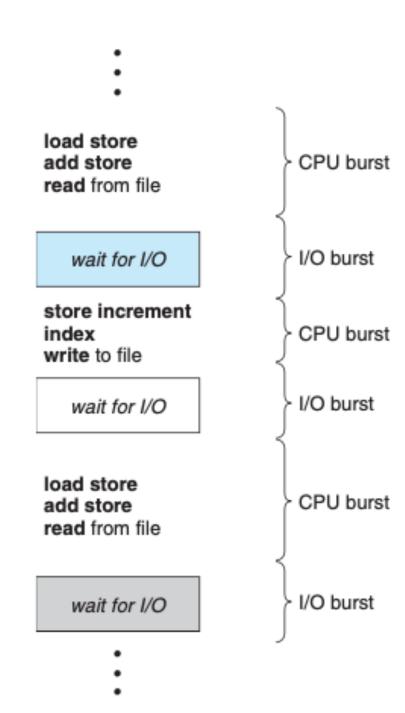
Processes in Ready Queue





Basic Concepts

 CPU-I/O Burst Cycle: Process execution consists of a cycle of CPU execution and I/O wait



CPU Scheduler

- Selects from among the processes in ready queue, and allocates the CPU to one of them
 - FIFO queue
 - Priority queue
 - Tree
 - Unordered linked-list
- CPU scheduling decisions may take place when a process:
 - 1. Switches from running to waiting state (I/O request)
 - 2. Switches from running to ready state (e.g. when interrupt occurs)
 - 3. Switches from waiting to ready (e.g. at completion of I/O)
 - 4. Terminates
- Scheduling under 1 and 4 is *nonpreemptive*
- All other scheduling is *preemptive*
 - Consider access to shared data
 - Consider preemption while in kernel mode
 - Consider interrupts occurring during crucial OS activities

Scheduling Criteria

- **CPU utilization** keep the CPU as busy as possible
- Throughput # of processes that complete their execution per time unit

• Turnaround time

- amount of time to execute a particular process
- -- the interval from the time of submission of a process to the time of the completion.
- -- sum of the periods spent waiting to get into memory, waiting in the ready queue, executing on the CPU, doing I/O
- Waiting time amount of time a process has been waiting in the ready queue
- **Response time** amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)

Scheduling Algorithm Optimization Criteria

- Max CPU utilization
- Max throughput
- Min turnaround time
- Min waiting time
- Min response time

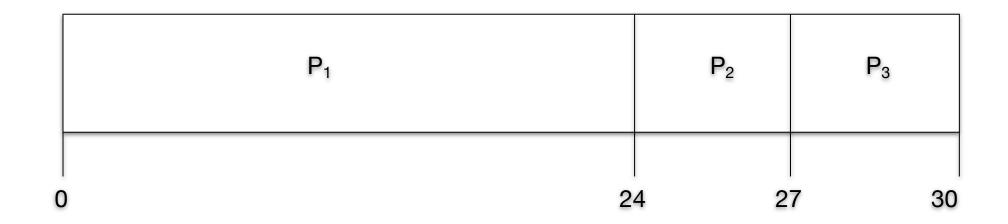
OPERATING SYSTEMS

CPU Scheduling Algorithms - First Come First Serve (FCFS)

First-Come, First-Served (FCFS) Scheduling

<u>Process</u>	Burst Time
P_1	24
P_2	3
P_3	3

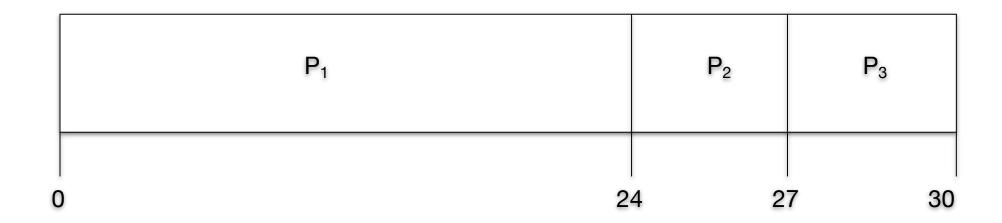
• Suppose that the processes arrive in the order: P_1 , P_2 , P_3 The Gantt Chart for the schedule is:



First-Come, First-Served (FCFS) Scheduling

<u>Process</u>	Burst Time
P_1	24
P_2	3
P_3	3

• Suppose that the processes arrive in the order: P_1 , P_2 , P_3 The Gantt Chart for the schedule is:



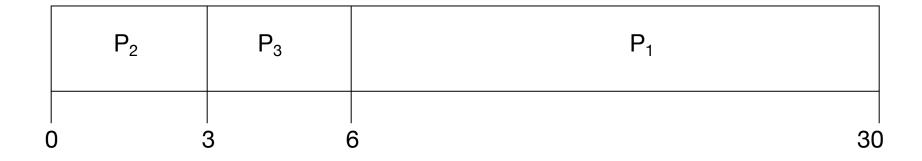
- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- Average waiting time: (0 + 24 + 27)/3 = 17
- Turnaround time $P_1 = 24$; $P_2 = 27$; $P_3 = 30$

FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order:

$$P_2, P_3, P_1$$

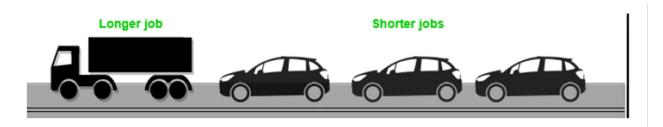
• The Gantt chart for the schedule is:



Waiting time for $P_1 = 6$; $P_2 = 0$, $P_3 = 3$ Average waiting time: (6 + 0 + 3)/3 = 3Much better than previous case

Convoy effect - short process behind long process.

Consider one CPU-bound and many I/O-bound processes



OPERATING SYSTEMS

CPU Scheduling Algorithms - Shortest Job First (SJF)

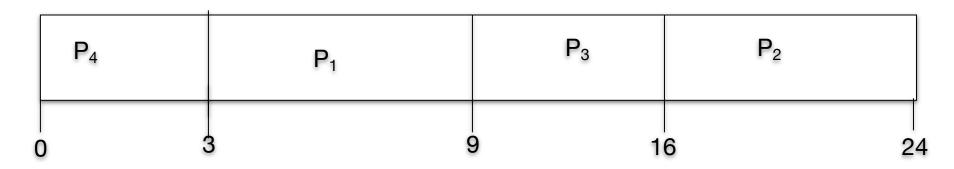
Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time
- Two schemes:
 - Non-preemptive once CPU given to the process it cannot be preempted until completes its CPU burst
 - o **preemptive** if a new process arrives with CPU burst length less than remaining time of current executing process, preempt. This scheme is known as the Shortest-Remaining-Time-First (SRTF)
- SJF is optimal gives minimum average waiting time for a given set of processes

Example of SJF

ProcessBurst Time P_1 6 P_2 8 P_3 7 P_4 3

• SJF scheduling chart

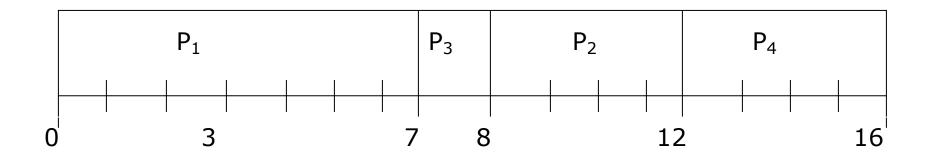


Waiting time for $P_1 = 3$; $P_2 = 16$; $P_3 = 9$; P4 = 0Average waiting time: (3 + 16 + 9 + 0)/4 = 7 ms

Example of Non-Preemptive SJF

<u>Process</u>	<u> Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

• SJF (non-preemptive)



Waiting time for
$$P_1 = 0$$
; $P_2 = (8 - 2) = 6$; $P_3 = (7 - 4) = 3$; $P_4 = (12 - 5) = 7$
Average waiting time: $(0 + 6 + 3 + 7)/4 = 4$ ms

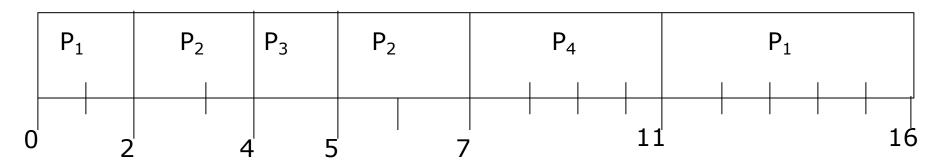
Average turnaround time: =
$$(7 + (12 - 2) + (8 - 4) + (16 - 5)) / 4 =$$

= $(7 + 10 + 4 + 11)/4 = 8ms$

Example of Preemptive SJF

<u>Process</u>	Arrival Time	Burst Time
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

• SJF (preemptive)



Waiting time for
$$P_1=0+(11-2)=9$$
; $P_2=(2-2)+(5-4)=1$; $P_3=(4-4)=0$; $P_4=(7-5)=2$
Average waiting time: $(9+1+0+2)/4=3$ ms

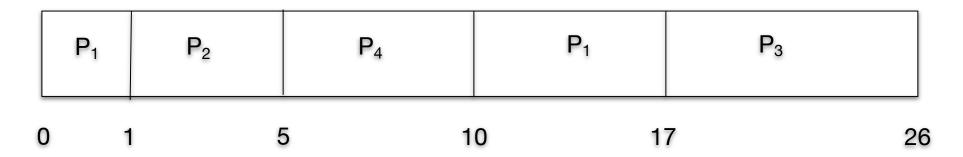
Average turnaround time: = ((16 - 0) + (7 - 2) + (5 - 4) + (11 - 5)) / 4 == (16 + 5 + 1 + 6)/4 = 7ms

Example of Shortest-remaining-time-first

• Now we add the concepts of varying arrival times and preemption to the analysis

<u>Process</u>	<u>Arrival Time</u>	Burst Time
P1	0	8
P2	1	4
P3	2	9
P4	3	5

Preemptive SJF Gantt Chart



Waiting time for
$$P_1 = 0 + (10 - 1) = 9$$
; $P_2 = (1 - 1) = 0$; $P_3 = (17 - 2) = 15$; $P_4 = (5 - 3) = 2$
Average waiting time: $(9 + 0 + 15 + 2)/4 = 6.5$ ms

Average turnaround time: =
$$((17 - 0) + (5 - 1) + (26 - 2) + (10 - 3)) / 4 =$$

= $(17 + 4 + 24 + 7)/4 = 13ms$

OPERATING SYSTEMS

CPU Scheduling Algorithms - Priority Scheduling

Priority Scheduling

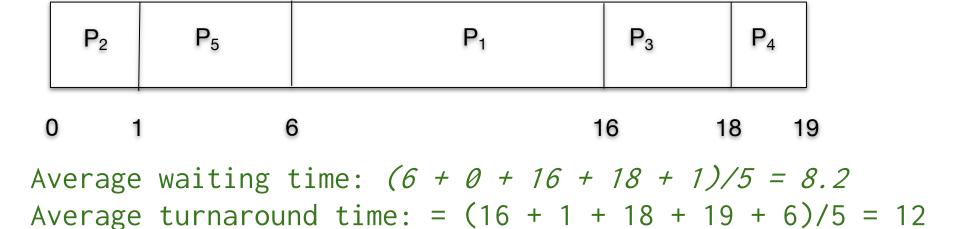
- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer ≡ highest priority)
 - Preemptive
 - Nonpreemptive
- SJF is priority scheduling where priority is the inverse of predicted next
 CPU burst time
- Priority can be defined either internally or externally.
 - Factors for internal priority assignment:
 - Time limit, memory requirements, the number or open files etc.
 - Factors for external priority assignment:
 - Importance of the process, the type and amount of funds being paid for computer use, department sponsoring works etc.

Example of Priority Scheduling

Non Preemptive:

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2

• Priority scheduling Gantt Chart



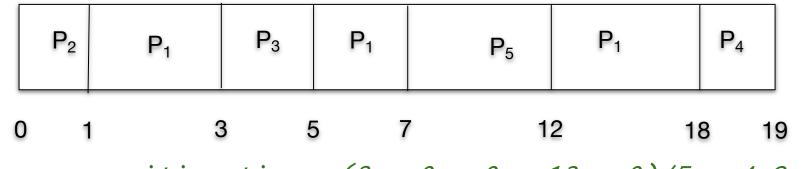
- Problem ≡ **Starvation** low priority processes may never execute
- Solution \equiv Aging as time progresses increase the priority of the process

Example of Priority Scheduling

Preemptive:

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	0	10	4
P_2	0	1	1
P_3	3	2	3
P_4	5	1	5
P_5	7	5	2

Priority scheduling Gantt Chart



Average waiting time: (8 + 0 + 0 + 13 + 0)/5 = 4.2Average turnaround time: = (18 + 1 + 2 + 14 + 5)/5 = 8

- Problem ≡ **Starvation** low priority processes may never execute
- Solution \equiv **Aging** as time progresses increase the priority of the process

OPERATING SYSTEMS

CPU Scheduling Algorithms - Round Robin (RR)

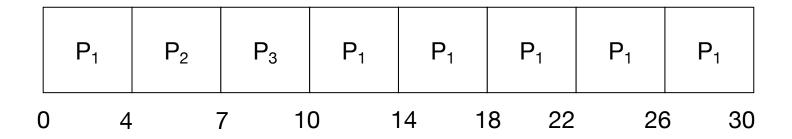
Round Robin (RR)

- Each process gets a small unit of CPU time (**time quantum** q), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are n processes in the ready queue and the time quantum is q, then each process gets 1/n of the CPU time in chunks of at most q time units at once. No process waits more than (n-1)q time units.
- Timer interrupts every quantum to schedule next process
- Performance
 - $q \text{ large} \Rightarrow \text{FIFO}$
 - $q \text{ small} \Rightarrow q \text{ must be large with respect to context switch, otherwise}$ overhead is too high

Example of RR with Time Quantum = 4

<u>Process</u>	Burst Time
P_1	24
P_2	3
P_3	3

• The Gantt chart is:



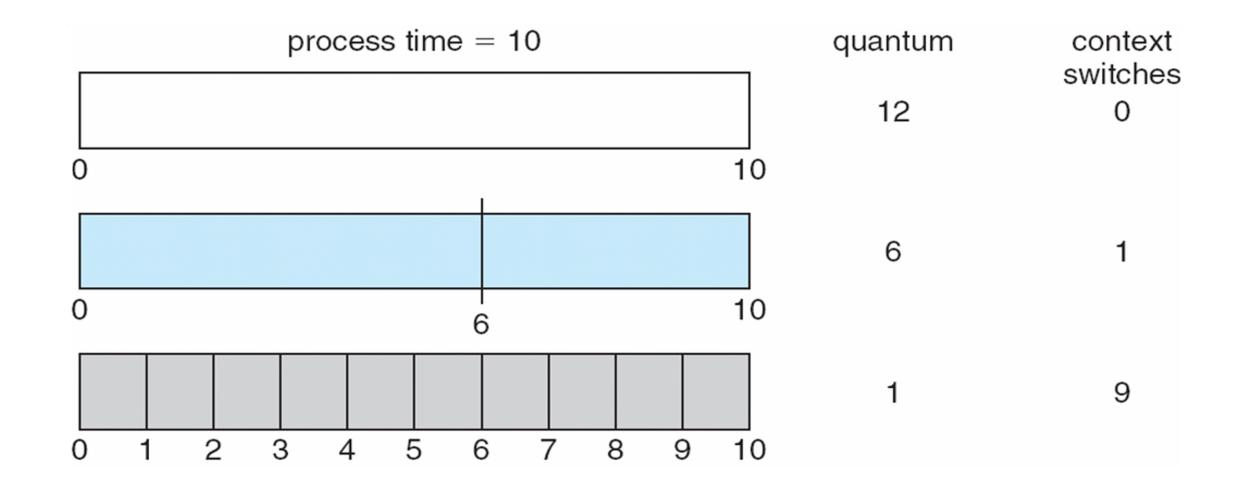
- Average waiting time is 17/3 = 5.66 milisecond
- Typically, higher average turnaround than SJF, but better response
- quantum should be large compared to context switch time
- Quantum usually 10ms to 100ms, context switch < 10 usec

Example of RR with Time Quantum = 2

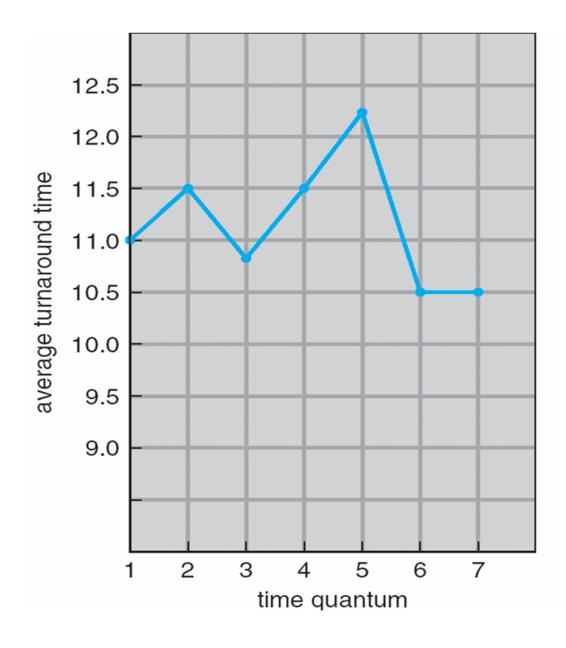
<u>Process</u>	<u>Arrival T</u>	<u> </u>	<u>Bı</u>	<u>urst Tin</u>	<u>ne</u>				
P1	0			5					
<i>P2</i>	1			3					
Р3	2			1					
<i>P4</i>	3			2					
<i>P5</i>	4			3					
	0	2 4	5	5 7	, (9 1	1 1	12	13 14
	P1	P2	P3	P1	P4	P5	P2	P1	P5

- Average waiting time is (8 + 8 + 2 + 4 + 7) / 5 = 29 / 5 = 5.8
- Average waiting time is (13 + 11 + 3 + 6 + 10) / 5 = 43 / 5 = 8.6

Time Quantum and Context Switch Time



Turnaround Time Varies With The Time Quantum



process	time
P_1	6
P_2	3
P_3	1
P_4	7

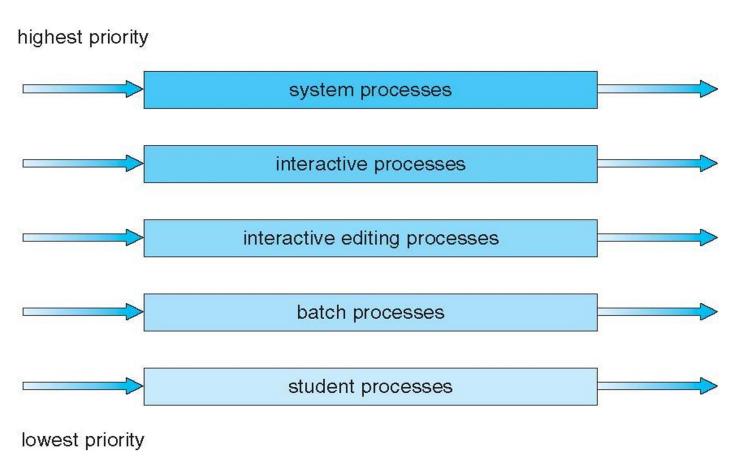
80% of CPU bursts should be shorter than quantum

OPERATING SYSTEMS

CPU Scheduling Multilevel Queue, Multilevel Feedback Queue

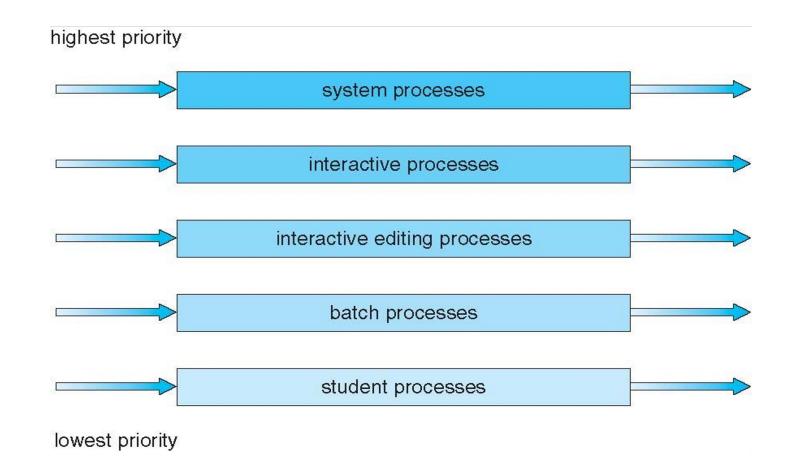
Multilevel Queue

- Another class of scheduling algorithm needs- in which processes are classified into different groups, e.g.:
 - foreground (interactive) processes
 - background (batch) processes
- They have different response time requirements-so different scheduling needs.
- Foreground processes may have priority over background processes.
- A multilevel queue-scheduling algorithm partitions the ready queue into several separate queues



Multilevel Queue

- Each queue has its own scheduling algorithm:
 - Foreground queue scheduled by RR algorithm
 - Background queue scheduled by FCFS algorithm
- Scheduling must be done between the queues:
 - Fixed priority preemptive scheduling; (i.e., serve all from foreground then from background). Possibility of starvation.
 - or Time slice each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e., foreground queue can be given 80% of the CPU time for RR-scheduling among its processes, while 20% to background in FCFS manner.



Multilevel Feedback Queue scheduling

- Multilevel Feedback Queue scheduling, allows a process to move between queues.
- If a process uses too much CPU time, it will be moved to a lower priority queue.
- Similarly, a process that waits too long in a lower-priority queue may me moved to a higher-priority queue.
- Multilevel-feedback-queue scheduler defined by the following parameters:
 - number of queues
 - scheduling algorithms for each queue
 - method used to determine when to upgrade a process
 - method used to determine when to demote a process
 - method used to determine which queue a process will enter when that process needs service

Example of Multilevel Feedback Queue

• Three queues: (can see the figure in next slide)

- Q_0 RR with time quantum 8 milliseconds
- Q_1 RR time quantum 16 milliseconds
- \circ $Q_2 FCFS$

Scheduling

- A new job enters queue Q_0 which is served for RR
 - When it gains CPU, job receives 8 milliseconds
 - If it does not finish in 8 milliseconds, job is moved to queue Q_1
- At Q₁ job is again served RR and receives 16 additional milliseconds
 - If it still does not complete, it is preempted and moved to queue Q_2

