### OPERATING SYSTEMS Process Synchronization

#### **BACKGROUND**

- Processes can execute concurrently or in parallel
- CPU scheduler switches rapidly between processes to provide concurrent execution
- A process may be interrupted at any point in its instruction stream
- Parallel execution, in which two instruction streams execute simultaneously on separate processing cores
- We will explain how concurrent or parallel execution can contribute to issues involving the integrity of data shared by several processes

#### PRODUCER-CONSUMER PROBLEM: Producer

```
item next_produced;
while (true) {
    /* produce an item in next_produced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```

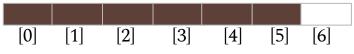
in: next free position in buffer
out: first full position in buffer

Both initialized with 0.

in = 0 out = 0

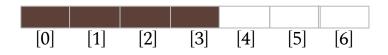
Here, BUFFER\_SIZE = 7
When buffer is full,

in = 6, out = 0



When buffer is not full,

$$in = 4$$
, out = 0



#### PRODUCER-CONSUMER PROBLEM: Consumer

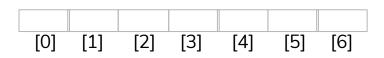
```
item next_consumed;
while (true) {
    while (in == out)
        ; /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    /* consume the item in next_consumed */
}
```

in: next free position in buffer
out: first full position in buffer

Both initialized with 0.

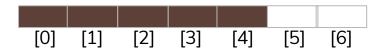
Here, BUFFER\_SIZE = 7

When buffer is empty, in = 0, out = 0



When buffer is not empty,

$$in = 5$$
, out = 0



#### PRODUCER-CONSUMER PROBLEM

- Modify the algorithm to remedy this deficiency add an integer variable *counter*, initialized to 0
- *counter* is incremented every time we add a new item to the buffer
- decremented every time we remove one item from the buffer

```
while (true) {
    /* produce an item in next_produced */
    while (counter == BUFFER_SIZE)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
    counter++;
}
```

```
while (true) {
    while (counter == 0)
      ; /* do nothing */

    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    counter--;

/* consume the item in next_consumed */
}
```

#### DATA INTEGRITY PROBLEM

```
register_1 = counterregister_2 = counterregister_1 = register_1 + 1register_2 = register_2 - 1counter = register_1counter = register_2
```

- "counter++" and "counter--" in machine language is like in the above.
- register1 and register2 is local CPU registers.
- Concurrent execution of "counter++" and "counter--" and allowing them to manipulate the counter variable create incorrect state.

```
producer
                           register_1 = counter
                                                       \{register_1 = 5\}
                execute
   producer execute
                           register_1 = register_1 + 1 \quad \{register_1 = 6\}
                           register_2 = counter
                                                  \{register_2 = 5\}
    consumer execute
                           register_2 = register_2 - 1 \quad \{register_2 = 4\}
   consumer execute
                           counter = register_1
T<sub>4</sub>: producer execute
                                                       \{counter = 6\}
                           counter = register_2
                                                       \{counter = 4\}
    consumer
                execute
```

#### RACE CONDITION

- Several process access and manipulate the same data concurrently
- Outcome of the execution depends on the particular order in which the access takes place
- To guard against this condition
  - Ensure that only one process at a time can manipulate the counter variable (shared data)
  - The processes should be synchronized

### OPERATING SYSTEMS Critical Section

#### **CRITICAL SECTION**

- Consider a system consisting of n processes  $\{P_0, P_1, \dots, P_{n-1}\}$ .
- *Critical Section:* segment of code of each process, which may change common variables, update a table, write a file and so on.
- While one process execute its critical section, no other process can execute their own critical section.
- Entry Section: section of code implementing critical section execution request
- Exit Section: section of code exiting from critical section
- Remainder section: Remaining code of the program.

#### REQUIREMENTS OF SOLUTION TO CRITICAL SECTION PROBLEM

1	Mutual exclusion:
	☐ If a process is executing its critical section, no other process can be executing in their critical sections.
2	Progress:
	☐ No process is executing in its critical section
	☐ Some process wish to enter their critical sections
	☐ Only those, who are not executing in their remainder section can participate in deciding which will enter the CS.
	☐ This selection cannot be postponed indefinitely.
3	Bounded waiting:
	☐ Bound or Limit on number of times other process can enter their CS after a process has made request to enter its CS and the request is granted

#### CRITICAL SECTIONS IN OPERATING SYSTEMS

Two general approach to handle CS in Operating System -

- 1. Preemptive kernel: allows a process to be preempted while it is running in kernel mode
- 2. Non-preemptive kernel: a kernel-mode process will run until it exits kernel mode, blocks, or voluntarily yields control of the CPU
- Non-preemptive kernel is free from race condition
- Preemptive kernel must be carefully designed to ensure that shared kernel data are free from race condition

# OPERATING SYSTEMS Peterson's solution for Critical Section Problem

#### SOFTWARE-BASED SOLUTION TO THE CRITICAL SECTION PROBLEM

- Known as "Perterson's Solution"
- Restricted to two processes that alternate execution between their critical sections and remainder sections
- Peterson's solution requires the two processes to share two data items:

int turn;

boolean flag[2];

- turn: indicates whose turn it is to enter its critical section
- flag: an array used to indicate if a process is ready to enter its critical section.

#### PETERSON'S SOLUTION

```
do {
     flag[i] = true;
     turn = j;
     while (flag[j] && turn == j);
         critical section
     flag[i] = false;
         remainder section
} while (true);
   The structure of process P_i in Peterson's solution.
```

- Process **Pi** first sets **flag[i]** to be **true** and then sets **turn** to the value **j**, so that if the other process wants to enter its CS, it can do so.
- If both try to enter at the same time, turn will be both *i* and *j* at the same time, but, only one of these will last.
- The eventual value of *turn* determines which process will enter its critical section.

It may not work correctly on modern computer architecture as they perform basic machine-language instructions such as load and store.

#### PETERSON'S SOLUTION

```
do {
do {
                                                             flag [ j ] = true;
    flag [i] = true;
                                                             turn = i;
    turn = j;
                                                              while ( flag [ i ] && turn == [ i ] );
     while ( flag [j] && turn == [j]);
                                                               critical section
      critical section
                                                             flag [ j ] = false;
    flag [i] = false;
                                                               remainder section
      remainder section
                                                               while (TRUE);
     } while (TRUE):
```

- Process Pi first sets flag[i] to be true and then sets turn to the value j, so that if the other process wants to enter its CS, it can do so.
  If both try to enter at the same
- If both try to enter at the same time, turn will be both *i* and *j* at the same time, but, only one of these will last.
- The eventual value of turn determines which process will enter its critical section.

It may not work correctly on modern computer architecture as they perform basic machine-language instructions such as load and store.

Process Pi

Process P<sub>j</sub>

#### Example

- Each Statement takes 2ms to execute, Process 1 gets executed first
- Context Switch will occur after 6ms
- Critical section contains 4 statements
- Remainder section contains 2 statements
- turn=0

} while (true);

• Flag[0] = FALSE, flag[1] = TRUE

```
flag[i] = true;
turn = j;
while (flag[j] && turn == j);

critical section

flag[i] = false;

remainder section
```

The structure of process  $P_i$  in Peterson's solution.

```
Process 0 (i = 0, j = 1)
                                     Process 1 (i = 1, j = 0)
                                    flag[1] = TRUE
                                    turn = 0
                                    While loop condition
flag[0] = TRUE
turn = 1
Stuck in while loop
                                     CS1
                                     CS2
                                    CS3
Stuck in while loop
                                     CS4
                                    flag[1] = FALSE
                                     RS1
While loop condition
CS1
CS2
                                     RS2
CS3
CS4
flag[0] = FALSE
RS1
RS2
```

# OPERATING SYSTEMS Hardware based solution for Critical Section Problem

#### HARDWARE-BASED SOLUTION TO THE CRITICAL SECTION PROBLEM

- More solutions to the critical-section problem using techniques ranging from hardware to software-based APIs
- These solutions are based on the premise of **locking** protecting critical regions through the use of locks.
- In a single-processor environment CS problem can be solved by preventing interrupts from occurring while a shared variable is being modified.
- For multiprocessor environment, we need different measures.
- Modern computer systems allow to test and modify the content of a word or to swap the contents of two words atomically which is uninterruptable unit. We can use <a href="test\_and\_set()">test\_and\_set()</a> and <a href="compare\_and\_swap()">compare\_and\_swap()</a> instructions.

#### TEST\_AND\_SET()

Executed atomically Mutual exclusion can be implemented by initializing a Boolean variable lock to false

```
boolean test_and_set(boolean *target) {
  boolean rv = *target;
  *target = true;

return rv;
}
```

```
do {
    while (test_and_set(&lock))
      ; /* do nothing */

    /* critical section */

    lock = false;

    /* remainder section */
} while (true);
```

#### TEST\_AND\_SET()

Executed atomically

Mutual exclusion can be implemented by initializing a Boolean variable lock to false

```
boolean TestAndSet (boolean *target) {
                         boolean rv = *target;
                                                                             Atomic
                         *target = TRUE;
                                                                           Operation
                         return rv;
                    The definition of the TestAndSet () instruction
                                                           do {
do {
                                                           while (TestAndSet (&lock));
while (TestAndSet (&lock));
// do nothing
                                                           // do nothing
                                                           // critical section
// critical section
                                                           lock = FALSE;
lock = FALSE;
                                                           // remainder section
// remainder section
                                                           } while (TRUE);
} while (TRUE);
```

#### COMPARE\_AND\_SWAP()

- Mutual exclusion can be achieved by declaring a global variable lock and initializing it to 0
- First process that invokes this instruction will set *lock* to 1 and no other process can execute CS until this process updates it to 0 after CS execution.

```
int compare_and_swap(int *value, int expected, int new_value) {
  int temp = *value;

if (*value == expected)
    *value = new_value;

  return temp;
}

do {
  while (compare_and_swap(&lock, 0, 1) != 0)
    ; /* do nothing */
    /* critical section */

  lock = 0;

  /* remainder section */
} while (true);
```

#### MUTEX LOCKS

- Operating-systems designers build software tools to solve CS problem.
- Simplest of these tools is "Mutex Lock" (Mutex = Mutual Exclusion)
- A process must acquire the lock before entering CS [ acquire() function ]
- A process must release the lock after exiting the CS [ release() function ]
- Mutex lock has a variable, available which indicates if the lock is available.

```
acquire() {
    while (!available)
    ; /* busy wait */
    available = false;;
}

release() {
    available = true;
} while (true);
remainder section
```

Figure 5.8 Solution to the critical-section problem using mutex locks.

### OPERATING SYSTEMS Semaphore

#### **SEMAPHORE**

- A semaphore S is an integer variable
- Is accessed only through two standard atomic operations: wait() and signal().
- When one process modifies the semaphore value, no other process can simultaneously modify that same semaphore value.
- In case of wait(S), the testing of the integer value of S ( $S \le 0$ ), as well as its possible modification (S--), must be executed without interruption, i.e., **this operations are atomic**.

#### Types of Semaphores

- □Counting Semaphore: The value can range over an unrestricted domain.
  - ☐ Used to control access to a given resource consisting of finite number of instances
  - □ Solves various synchronization problems.
- □Binary Semaphore: The value can range only between 0 and 1. This behaves similar to Mutex Lock.

#### **Counting Semaphore**

- initialized to the number of resources available, S = n
- Each process that wishes to use a resource performs a wait() operation

$$S = S - 1$$

• When a process releases a resource, it performs a signal() operation

$$S = S + 1$$

- When **S** becomes **0**, all resources are being used
- processes that wish to use a resource will block until **S**>**0**

#### Binary Semaphore - Synchronization

- P<sub>1</sub> has statement S<sub>1</sub>
- P<sub>2</sub> has statement S<sub>2</sub>

- We want to make sure that  $S_1$  executes before  $S_2$
- We can use a semaphore variable sync and initialize it to 0

```
      P1:
      P2:

      S1;
      wait(sync);

      signal(sync);
      S2;
```

#### Mutual Exclusion With Semaphores

- Binary Semaphores (mutex) can be used to solve CS problem.
- A semaphore variable (say mutex) can be shared by n processes and initialized to 1.
- Each process is structured as follows:

```
do{
    wait (mutex);
    //critical section
    signal(mutex);
    //remainder section
}while (TRUE);
```

#### Deadlock & Starvation

- Two or more process can wait indefinitely for an event **DEADLOCK** !!!
- It occurs because two process depends on each other for causing an event in a specific manner

```
\begin{array}{cccc} P_0 & P_1 \\ \text{wait}(S); & \text{wait}(Q); \\ \text{wait}(Q); & \text{wait}(S); \\ & \ddots & & \ddots \\ & \dots & & \dots \\ & \text{signal}(S); & \text{signal}(Q); \\ \text{signal}(Q); & \text{signal}(S); \end{array}
```

- **Starvation**: Processes wait indefinitely within the semaphore
- Occurs if we remove processes from the list associated with a semaphore in LIFO (last-in, first-out) order

### OPERATING SYSTEMS Semaphore Implementation

#### SEMAPHORE IMPLEMENTATION

- Must guarantee that no two processes can execute the **wait**() and **signal**() on the same semaphore at the same time
- Thus, the implementation becomes the critical section problem where the wait and signal code are placed in the critical section
- Could now have busy waiting in critical section implementation
  - But implementation code is short
  - Little busy waiting if critical section rarely occupied
- Note that applications may spend lots of time in critical sections and therefore this is not a good solution

#### SEMAPHORE IMPLEMENTATION

- When a process executes the wait() operation and finds that the semaphore value is not positive, it must wait
- Rather than this busy waiting, the process can block itself which places it into a waiting queue associated with the semaphore
- State of the process is switched to the waiting state and control is transferred to CPU scheduler which selects another process to execute.
- It will be restarted when some other process executes a signal() operation
- Restarted by a wakeup() operation that changes it from waiting state to ready state.

#### SEMAPHORE IMPLEMENTATION

#### Definition of a semaphore:

```
typedef struct {
    int value;
    struct process *list;
} semaphore;
```

# Definition of wait(): wait(semaphore \*S) { S->value--; if (S->value < 0) { add this process to S->list; block(); }

```
Definition of signal():

signal(semaphore *S) {
    S->value++;
    if (S->value >= 0) {
        remove a process P
        from S->list;
        wakeup(P);
    }
}
```